

## ▼ Data Ingestion & Profiling

**Objective:** Collect and preprocess diverse text data from "Human" and "AI" sources to augment the training set. This notebook handles web scraping, API fetching, and parsing personal data dumps.

### Data Sources Processed:

- **Wikipedia (Human):** Scrapped articles across STEM, History, and Culture topics.
- **arXiv (Human):** Abstract summaries from ML, Stats, and Physics papers.
- **News Feeds (Human):** RSS articles from BBC, Reuters, and AP.
- **ChatGPT History (AI):** Personal conversations exported (with consent) from OpenAI (Label 1).

### Output:

- Generates `scraped_data_combined.csv` (saved to Drive), which will be merged with the larger Kaggle dataset in Notebook 2.

```
!pip -q install pandas numpy scikit-learn tqdm wikipedia feedparser beautifulsoup4 readability-lxml requests

Preparing metadata (setup.py) ... done
Preparing metadata (setup.py) ... done
                               81.5/81.5 kB 2.0 MB/s eta 0:00:00
Building wheel for wikipedia (setup.py) ... done
Building wheel for sgmlplib3k (setup.py) ... done
```

```
import os, re, json, time, hashlib
import numpy as np
import pandas as pd
from tqdm.notebook import tqdm
from sklearn.model_selection import train_test_split

import requests
from bs4 import BeautifulSoup
from readability import Document
import feedparser
import wikipedia

tqdm.pandas()

RANDOM_SEED = 42
np.random.seed(RANDOM_SEED)

TEXT_COL = "text"
LABEL_COL = "label"      # 0=human, 1=AI

def clean_text(t: str) -> str:
    if t is None:
        return ""
    t = re.sub(r"http\S+", " ", str(t))
    t = re.sub(r"\S+@\S+", " ", t)
    t = re.sub(r"\s+", " ", t).strip()
    return t
```

```

def stable_hash(s: str) -> str:
    return hashlib.sha1(s.encode("utf-8", errors="ignore")).hexdigest()

def chunk_text_words(text: str, chunk_words=200, overlap=40, min_words=40):
    words = re.findall(r"\S+", text)
    if len(words) < min_words:
        return []
    step = max(1, chunk_words - overlap)
    chunks = []
    for i in range(0, len(words), step):
        ch = " ".join(words[i:i+chunk_words]).strip()
        if len(ch.split()) >= min_words:
            chunks.append(ch)
    return chunks

def add_len_bins(df: pd.DataFrame):
    df["len_words"] = df[TEXT_COL].str.split().str.len()
    df["len_bin"] = pd.cut(df["len_words"], bins=[0,10,25,50,100,200,400,1000,10_000],
                           labels=False, include_lowest=True)
    return df

def dedup_by_text(df: pd.DataFrame):
    df["text_hash"] = df[TEXT_COL].map(stable_hash)
    df = df.drop_duplicates(subset=["text_hash"]).drop(columns=["text_hash"])
    return df

```

```

from google.colab import drive
drive.mount("/content/drive")

```

Mounted at /content/drive

```

import pandas as pd
import re

```

```

CHATGPT_JSON = "/content/drive/MyDrive/conversations.json" # upload this

with open(CHATGPT_JSON, "r", encoding="utf-8") as f:
    convos = json.load(f)

rows = []
for convo in convos:
    convo_id = convo.get("id", "unknown_convo")
    mapping = convo.get("mapping", {})
    for node in mapping.values():
        msg = node.get("message")
        if not msg:
            continue
        role = msg.get("author", {}).get("role")
        parts = (msg.get("content", {}) or {}).get("parts", []) or []
        if role == "assistant" and parts:
            text = clean_text(" ".join([p for p in parts if isinstance(p, str)]))
            if len(text.split()) >= 20:
                rows.append({
                    "doc_id": f"chatgpt_{convo_id}",

```

```

        "source": "chatgpt",
        "text": text,
        "label": 1
    })

chatgpt_df = pd.DataFrame(rows)
print("ChatGPT raw samples:", chatgpt_df.shape)
chatgpt_df.head()

```

ChatGPT raw samples: (4178, 4)

	doc_id	source	text	label	
0	chatgpt_69351ea1-82b0-8325-8e60-8b957ba06590	chatgpt	Short answer: treat this like *two* binary cla...	1	
1	chatgpt_69351ea1-82b0-8325-8e60-8b957ba06590	chatgpt	Text. 100%. Given your background and what you...	1	
2	chatgpt_69351ea1-82b0-8325-8e60-8b957ba06590	chatgpt	Perfect — here is a **clean, actionable list**...	1	
3	chatgpt_69351ea1-82b0-8325-8e60-8b957ba06590	chatgpt	Nice, that Kaggle dataset is a good choice. Le...	1	
4	chatgpt_69351ea1-82b0-8325-8e60-8b957ba06590	chatgpt	Yes — you **already have a proper validation s...	1	

Next steps: [Generate code with chatgpt\\_df](#) [New interactive sheet](#)

```

chat_chunks = []
for _, r in chatgpt_df.iterrows():
    for ch in chunk_text_words(r["text"], chunk_words=180, overlap=40, min_words=40):
        chat_chunks.append({"doc_id": r["doc_id"], "source": r["source"], "text": ch, "label": r["label"]})

chatgpt_df = pd.DataFrame(chat_chunks)
print("ChatGPT chunks:", chatgpt_df.shape)

#So that long outputs don't dominate

```

ChatGPT chunks: (12179, 4)

```

WIKI_TOPICS = [
    # --- Core AI / CS / Tech ---
    "Artificial intelligence",
    "Machine learning",
    "Deep learning",
    "Natural language processing",
    "Computer vision",
    "Algorithms",
    "Data science",
    "Cloud computing",
    "Cybersecurity",
    "Blockchain",

    # --- Physical & Life Sciences ---
    "Quantum mechanics",
    "Relativity",
    "Thermodynamics",
    "Biochemistry",
    "Genetics",
    "Molecular biology",
    "Neuroscience",
    "Climate change",

```

```
"Evolution",
"Epidemiology",

# --- Mathematics & Logic ---
"Calculus",
"Linear algebra",
"Probability theory",
"Statistics",
"Graph theory",
"Number theory",
"Mathematical logic",

# --- Law, Policy & Society ---
"Indian Contract Act",
"Constitution of India",
"International law",
"Human rights",
"Intellectual property",
"Privacy law",
"Cyber law",
"Regulation of artificial intelligence",

# --- Economics, Business & Finance ---
"Macroeconomics",
"Microeconomics",
"Behavioral economics",
"Game theory",
"Financial markets",
"Inflation",
"Globalization",

# --- History & Politics ---
"World War I",
"World War II",
"Cold War",
"French Revolution",
"Indian independence movement",
"United States presidential election",
"Geopolitics",

# --- Literature, Language & Arts ---
"William Shakespeare",
"English literature",
"Poetry",
"Drama",
"Literary criticism",
"Rhetoric",
"Comparative literature",

# --- Philosophy & Social Thought ---
"Philosophy",
"Ethics",
"Epistemology",
"Metaphysics",
"Existentialism",
"Philosophy of science",

# --- Media, Sports & Culture ---
```

```

    "Formula One",
    "Olympic Games",
    "Association football",
    "Cricket",
    "Film theory",
    "Music theory",
    "Popular culture",

    # --- Current / General ---
    "Current events",
    "Global health",
    "Sustainability",
    "Energy transition",
    "Artificial general intelligence"
]

wiki_docs = []
for topic in WIKI_TOPICS:
    try:
        page = wikipedia.page(topic, auto_suggest=False)
        text = clean_text(page.content)
        wiki_docs.append({
            "doc_id": f"wiki_{page.pageid}",
            "source": "wikipedia",
            "text": text,
            "label": 0
        })
    except Exception as e:
        print("Skip wiki:", topic, "|", str(e)[:120])

wiki_docs = pd.DataFrame(wiki_docs)

wiki_chunks = []
for _, r in wiki_docs.iterrows():
    for ch in chunk_text_words(r["text"], chunk_words=220, overlap=50, min_words=60):
        wiki_chunks.append({"doc_id": r["doc_id"], "source": r["source"], "text": ch, "label": r["label"]})

wiki_df = pd.DataFrame(wiki_chunks)
print("Wikipedia chunks:", wiki_df.shape)

```

```

/usr/local/lib/python3.12/dist-packages/wikipedia/wikipedia.py:389: GuessedAtParserWarning: No parser was explicitly specified, so I'm using the best available
The code that caused this warning is on line 389 of the file /usr/local/lib/python3.12/dist-packages/wikipedia/wikipedia.py. To get rid of this warning, pass th
lis = BeautifulSoup(html).find_all('li')
Skip wiki: Relativity | "Relativity" may refer to:
Galilean relativity
Numerical relativity
Principle of relativity
Theory of relativity
Genera
Wikipedia chunks: (3355, 4)

```

Start coding or [generate](#) with AI.

```

import time
import urllib.parse
import feedparser
import pandas as pd

# import re
# def clean_text(t: str) -> str:
#     t = re.sub(r"http\S+", " ", str(t))
#     t = re.sub(r"\S+@\S+", " ", t)
#     t = re.sub(r"\s+", " ", t).strip()
#     return t

ARXIV_QUERIES = [
    "cat:stat.ML",           # Stats ML
    "cat:math.ST",           # Statistics
    "cat:physics.optics",    # Physics
    "cat:q-bio.BM"           # Bioinformatics
]

MAX_RESULTS_PER_QUERY = 200      # per category
ARXIV_START = 0
MIN_WORDS = 60                  # keep decent-length abstracts
SLEEP_SECONDS = 3               # be polite to arXiv

paper_rows = []
seen_ids = set()

for q in ARXIV_QUERIES:
    encoded_q = urllib.parse.quote(q)
    url = (
        "http://export.arxiv.org/api/query"
        f"?search_query={encoded_q}"
        f"&start={ARXIV_START}"
        f"&max_results={MAX_RESULTS_PER_QUERY}"
    )

    feed = feedparser.parse(url)
    print(f"Query={q} | fetched entries={len(feed.entries)}")

    for entry in feed.entries:
        arxiv_id = entry.id.split("/")[-1]
        if arxiv_id in seen_ids:
            continue
        seen_ids.add(arxiv_id)

        title = clean_text(getattr(entry, "title", ""))
        abstract = clean_text(getattr(entry, "summary", ""))

        text = f"{title}. {abstract}".strip()
        if len(text.split()) < MIN_WORDS:
            continue

        paper_rows.append({
            "doc_id": f"arxiv_{arxiv_id}",
            "source": "arxiv",
            "text": text,
            "label": 0
        })

```

```
time.sleep(SLEEP_SECONDS)

papers_df = pd.DataFrame(paper_rows).drop_duplicates(subset=["doc_id"]).reset_index(drop=True)

print("\nDone")
print("Total arXiv samples:", len(papers_df))
print(papers_df.head(3))

Query=cat:stat.ML | fetched entries=200
Query=cat:math.ST | fetched entries=200
Query=cat:physics.optics | fetched entries=200
Query=cat:q-bio.BM | fetched entries=200

Done
Total arXiv samples: 780
   doc_id source \
0  arxiv_2012.12056v1  arxiv
1  arxiv_2012.13115v1  arxiv
2  arxiv_2012.13190v2  arxiv

          text  label
0  Data Assimilation in the Latent Space of a Neu...      0
1  Upper Confidence Bounds for Combining Stochast...      0
2  QUACKIE: A NLP Classification Task With Ground...      0
```

```
paper_chunks = []
for _, r in papers_df.iterrows():
    for ch in chunk_text_words(r["text"], chunk_words=200, overlap=40, min_words=60):
        paper_chunks.append({"doc_id": r["doc_id"], "source": r["source"], "text": ch, "label": r["label"]})

papers_df = pd.DataFrame(paper_chunks)
print("arXiv chunks:", papers_df.shape)
```

```
#So long outputs don't dominate
```

```
arXiv chunks: (893, 4)
```

```
NEWS_RSS = [
    # --- BBC ---
    "https://feeds.bbci.co.uk/news/world/rss.xml",
    "https://feeds.bbci.co.uk/news/technology/rss.xml",
    "https://feeds.bbci.co.uk/news/business/rss.xml",
    "https://feeds.bbci.co.uk/news/science_and_environment/rss.xml",

    # --- NPR (US, high editorial quality) ---
    "https://www.npr.org/rss/rss.php?id=1001",    # News
    "https://www.npr.org/rss/rss.php?id=1019",    # Technology
    "https://www.npr.org/rss/rss.php?id=1007",    # Science
    "https://www.npr.org/rss/rss.php?id=1014",    # World

    # --- Reuters (excellent for neutral tone) ---
    "https://feeds.reuters.com/reuters/worldNews",
    "https://feeds.reuters.com/reuters/businessNews",
    "https://feeds.reuters.com/reuters/technologyNews",
    "https://feeds.reuters.com/reuters/scienceNews",

    # --- Associated Press ---
    "https://feeds.ap.org/rss/APWorldNews",
    "https://feeds.ap.org/rss/APBusinessNews",
    "https://feeds.ap.org/rss/APTechnologyNews",
    "https://feeds.ap.org/rss/APScienceNews"
]
```

```

"https://apnews.com/rss",
"https://apnews.com/hub/technology/rss",
"https://apnews.com/hub/science/rss",

# --- The Guardian ---
"https://www.theguardian.com/world/rss",
"https://www.theguardian.com/technology/rss",
"https://www.theguardian.com/science/rss",
"https://www.theguardian.com/business/rss",

# --- Financial / Economics ---
"https://www.ft.com/rss/home",           # Financial Times (some paywall, still usable)
"https://www.economist.com/rss"          # Economist (often truncated)
]

```

```

def fetch_article_text(url: str, timeout=10) -> str:
    try:
        r = requests.get(url, timeout=timeout, headers={"User-Agent": "Mozilla/5.0"})
        if r.status_code != 200:
            return ""
        doc = Document(r.text)
        html = doc.summary()
        soup = BeautifulSoup(html, "html.parser")
        text = clean_text(soup.get_text(" "))
        return text
    except Exception:
        return ""

news_rows = []
MAX_ARTICLES_PER_FEED = 50

```

```

for feed_url in NEWS_RSS:
    f = feedparser.parse(feed_url)
    for entry in f.entries[:MAX_ARTICLES_PER_FEED]:
        url = entry.get("link", "")
        title = clean_text(entry.get("title", ""))
        body = fetch_article_text(url)
        text = f"{title}. {body}".strip()
        if len(text.split()) >= 80:
            news_rows.append({
                "doc_id": f"news_{stable_hash(url)[:12]}",
                "source": "news",
                "text": text,
                "label": 0
            })
        time.sleep(0.2)  # be polite

```

```

news_df = pd.DataFrame(news_rows)
print("News samples:", news_df.shape)

```

```

News samples: (334, 4)

```

```

news_chunks = []
for _, r in news_df.iterrows():
    for ch in chunk_text_words(r["text"], chunk_words=220, overlap=50, min_words=80):
        news_chunks.append({"doc_id": r["doc_id"], "source": r["source"], "text": ch, "label": r["label"]})

```

```
news_df = pd.DataFrame(news_chunks)
print("News chunks:", news_df.shape)
```

```
News chunks: (1873, 4)
```

```
dfs_to_merge = []

if 'chatgpt_df' in locals(): dfs_to_merge.append(chatgpt_df)
if 'wiki_df' in locals(): dfs_to_merge.append(wiki_df)
if 'papers_df' in locals(): dfs_to_merge.append(papers_df)
if 'news_df' in locals(): dfs_to_merge.append(news_df)

if dfs_to_merge:
    scraped_combined = pd.concat(dfs_to_merge, ignore_index=True)
    scraped_combined["text"] = scraped_combined["text"].astype(str).str.strip()
    output_path = "/content/drive/MyDrive/scraped_data_combined.csv"
    print(f"Merging {len(dfs_to_merge)} data sources...")
    scraped_combined.to_csv(output_path, index=False)
    print("-" * 30)
    print(f"SUCCESS: Scrapped data merged and saved to: {output_path}")
    print(f"Total Combined Shape: {scraped_combined.shape}")
    print("Source Distribution:")
    print(scraped_combined["source"].value_counts())
else:
    print("ERROR: No dataframes found to merge. Please run the scraping cells first.")
```

```
Merging 4 data sources...
```

```
-----  
SUCCESS: Scrapped data merged and saved to: /content/drive/MyDrive/scraped_data_combined.csv  
Total Combined Shape: (18300, 4)  
Source Distribution:  
source  
chatgpt      12179  
wikipedia    3355  
news         1873  
arxiv        893  
Name: count, dtype: int64
```

```
#Final Completed
```