

04_lstm_sequence_model

December 14, 2025

1 LSTM Sequence Model

Objective: Train a Deep Learning sequence model (LSTM) to detect AI-generated text. Unlike the feature-based models in Notebook 3 which relied on manual counts (like sentence length), this model learns patterns directly from the raw sequence of words.

Key Components: 1. **Load Fixed Splits:** Import the pre-split data (`train.csv`, `val.csv`, `test.csv`) saved in Notebook 3 to ensure we are testing on the exact same “Collected” dataset. 2. **Preprocessing:** * **Tokenization:** Custom regex tokenizer to split text into words. * **Vocabulary:** Built *strictly* on the Training set to prevent data leakage. * **Sequence Handling:** Implements padding and “packed sequences” to handle variable-length text efficiently in PyTorch. 3. **Model Architecture:** * **Embedding Layer:** Converts words into dense vectors. * **Bi-LSTM:** Bidirectional Long Short-Term Memory layer to capture context from both past and future words. * **Classifier:** Fully connected layer with Dropout for regularization. 4. **Training Loop:** * Uses **AdamW** optimizer and **BCEWithLogitsLoss**. * Implements **Early Stopping** based on Validation F1-score to prevent overfitting. 5. **Final Evaluation:** Reports strict accuracy and F1 metrics on the held-out Test set (Scraped Data).

```
[16]: # === LOAD FIXED SPLITS (exported from baseline notebook) ===

from google.colab import drive
drive.mount("/content/drive")

import json
from pathlib import Path
import pandas as pd
import numpy as np

ART_DIR = Path("/content/drive/MyDrive/artifacts/data_splits_v1") # same ↴
    ↴ folder used in baseline

# --- load metadata ---
with open(ART_DIR / "meta.json") as f:
    meta = json.load(f)

fmt = meta["format"]
style_cols = meta["style_cols"]
```

```

# --- load datasets ---
if fmt == "parquet":
    train_df = pd.read_parquet(ART_DIR / "train_all.parquet")
    val_df   = pd.read_parquet(ART_DIR / "val_all.parquet")
    test_df  = pd.read_parquet(ART_DIR / "test_all.parquet")
else:
    train_df = pd.read_csv(ART_DIR / "train_all.csv")
    val_df   = pd.read_csv(ART_DIR / "val_all.csv")
    test_df  = pd.read_csv(ART_DIR / "test_all.csv")

# --- sanity checks (text + label + style columns) ---
required_cols = ["text", "label", "source"] + style_cols

for name, df in [("train", train_df), ("val", val_df), ("test", test_df)]:
    missing = [c for c in required_cols if c not in df.columns]
    if missing:
        raise ValueError(f"{name} split missing columns: {missing[:15]}{' ... '}" if len(missing) > 15 else '')

# --- labels as numpy arrays ---
y_train = train_df["label"].astype(int).values
y_val   = val_df["label"].astype(int).values
y_test  = test_df["label"].astype(int).values

print("Loaded splits from:", ART_DIR)
print("Format:", fmt)
print("Sizes:", len(train_df), len(val_df), len(test_df))
print("Label dist train:", np.bincount(y_train))
print("Label dist val:  ", np.bincount(y_val))
print("Label dist test: ", np.bincount(y_test))
print("Num stylometry features:", len(style_cols))

```

Drive already mounted at /content/drive; to attempt to forcibly remount, call
drive.mount("/content/drive", force_remount=True).

Loaded splits from: /content/drive/MyDrive/artifacts/data_splits_v1

Format: parquet

Sizes: 32615 5756 1629

Label dist train: [16677 15938]

Label dist val: [2943 2813]

Label dist test: [380 1249]

Num stylometry features: 33

[17]: train_df

		text	source	num_chars	\
0	Social Media has become a big part of our live...	kaggle	2189		
1	I am for the development OS these cars. Why, b...	kaggle	2177		

2	The Electoral College is a unique system in th...	kaggle	2039
3	Many of you people who believe in TLE paranorm...	kaggle	1224
4	Many schools set up summer projects to provide...	kaggle	3010

...
32610	Thomas Jefferson was a man of many ideas and a...	kaggle	2487
32611	Dear Principal, I do think it's fair to let on...	kaggle	1791
32612	Some people agree with making students partici...	kaggle	2706
32613	There are pros and cons to learning to compete...	kaggle	785
32614	Driverless cars, also known as autonomous cars...	kaggle	3038

	num_words	num_sentences	avg_sentence_length	type_token_ratio	\
0	411	26	15.807693	0.406326	
1	401	24	16.708334	0.438903	
2	341	16	21.312500	0.445748	
3	220	13	16.923077	0.531818	
4	520	27	19.259260	0.334615	
...
32610	437	25	17.480000	0.352403	
32611	357	15	23.799999	0.445378	
32612	460	24	19.166666	0.300000	
32613	134	6	22.333334	0.425373	
32614	474	24	19.750000	0.415612	

	unique_words	pct_punct	pct_upper	...	colon_ratio_emphasis	\
0	167	0.035633	0.023755	...	0.000000	
1	176	0.021589	0.022049	...	0.000459	
2	152	0.022070	0.020598	...	0.000000	
3	117	0.024510	0.059641	...	0.000000	
4	174	0.014286	0.031561	...	0.000000	
...
32610	154	0.015279	0.013269	...	0.000000	
32611	159	0.027917	0.026242	...	0.000000	
32612	138	0.015521	0.008869	...	0.000000	
32613	57	0.016561	0.014013	...	0.000000	
32614	197	0.019421	0.007900	...	0.000000	

	semicolon_ratio_emphasis	emphasis_punctuation_ratio	\
0	0.000000	0.000000	
1	0.000000	0.000459	
2	0.000000	0.000000	
3	0.000000	0.000000	
4	0.000332	0.000332	
...
32610	0.000000	0.000000	
32611	0.000558	0.000558	
32612	0.000000	0.000000	
32613	0.000000	0.000000	

32614	0.000000	0.000000		
	hedge_phrase_count	hedge_phrase_ratio	hedge_line_ratio	\
0	0	0.0	0.0	
1	0	0.0	0.0	
2	0	0.0	0.0	
3	0	0.0	0.0	
4	0	0.0	0.0	
...	
32610	0	0.0	0.0	
32611	0	0.0	0.0	
32612	0	0.0	0.0	
32613	0	0.0	0.0	
32614	1	1.0	1.0	
	generic_claim_ratio	vague_quantifier_ratio	fake_citation_ratio	label
0	0.0	1.0	0.0	1
1	0.0	0.0	0.0	0
2	0.0	1.0	0.0	1
3	0.0	1.0	0.0	0
4	0.0	1.0	0.0	0
...
32610	0.0	1.0	0.0	1
32611	0.0	1.0	0.0	0
32612	0.0	1.0	0.0	0
32613	0.0	0.0	0.0	1
32614	0.0	1.0	0.0	1

[32615 rows x 36 columns]

```
[3]: # =====
# LSTM text classifier
# TRAIN/VAL = Kaggle only
# TEST      = Scrapped only
#
# Uses PyTorch. Includes:
# - fast regex tokenization
# - vocab built from TRAIN only
# - truncation to max_len
# - padding + packed sequences
# - early stopping on VAL F1
# - per-epoch timing
#
# datasets:
#   train_df, val_df, test_df
#   y_train, y_val, y_test
#
```

```

# This code prints time/epoch + a simple projected total after epoch 1.
# =====

import re
import time
import math
import numpy as np
from collections import Counter

import torch
import torch.nn as nn
from torch.utils.data import Dataset, DataLoader

from sklearn.metrics import precision_recall_fscore_support, accuracy_score

# -----
# Config (tune these first)
# -----

SEED = 42
MAX_VOCAB = 50_000          # 30k-100k typical
MIN_FREQ = 2                 # drop very rare tokens, since we want to focus on
                             # statistical construct of the text
MAX_LEN = 384               # 256/384/512. Higher = slower but captures longer
                             # context
BATCH_SIZE = 32              # 32/64/128 depending on CPU & RAM
EMB_DIM = 192                # 128-256
HID_DIM = 192                # 128-256
NUM_LAYERS = 1                # 1-2 (2 slower)
BIDIR = True
DROPOUT = 0.2
LR = 2e-3
EPOCHS = 8
PATIENCE = 4                  # early stop if VAL F1 doesn't improve
CLIP = 1.0

DEVICE = torch.device("cuda" if torch.cuda.is_available() else "cpu")
print("Using device:", DEVICE)

# -----
# Reproducibility
# -----
torch.manual_seed(SEED)
np.random.seed(SEED)

# -----
# Tokenizer (fast, stable)
# -----

```

```

_tok = re.compile(r"[A-Za-z0-9]+(?:'[A-Za-z0-9]+)?|[^\\sA-Za-z0-9]")
def tokenize(text: str):
    return _tok.findall((text or "").lower())

# -----
# Build vocab from TRAIN only
# -----
def build_vocab(texts, max_vocab=MAX_VOCAB, min_freq=MIN_FREQ):
    counter = Counter()
    for t in texts:
        counter.update(tokenize(t))
    # Special tokens
    itos = ["<pad>", "<unk>"]
    # Keep most common above min_freq
    for tok, freq in counter.most_common():
        if freq < min_freq:
            break
        itos.append(tok)
        if len(itos) >= max_vocab:
            break
    stoi = {tok: i for i, tok in enumerate(itos)}
    return stoi, itos

train_texts = train_df["text"].astype(str).tolist()
val_texts = val_df["text"].astype(str).tolist()
test_texts = test_df["text"].astype(str).tolist()

stoi, itos = build_vocab(train_texts)
PAD_IDX = stoi["<pad>"]
UNK_IDX = stoi["<unk>"]

print(f"Vocab size: {len(itos)}, (PAD={PAD_IDX}, UNK={UNK_IDX})")

# -----
# Dataset + Collate
# -----
def encode(text: str, max_len=MAX_LEN):
    ids = [stoi.get(tok, UNK_IDX) for tok in tokenize(text)]
    if len(ids) > max_len:
        ids = ids[:max_len]
    return ids

class TextDataset(Dataset):
    def __init__(self, texts, labels):
        self.texts = texts
        self.labels = labels.astype(np.int64)
    def __len__(self):

```

```

        return len(self.texts)
    def __getitem__(self, i):
        return self.texts[i], self.labels[i]

    def collate_batch(batch):
        texts, labels = zip(*batch)
        seqs = [torch.tensor(encode(t), dtype=torch.long) for t in texts]
        lengths = torch.tensor([len(s) for s in seqs], dtype=torch.long)

        # pad to max length in batch
        maxl = int(lengths.max().item()) if len(lengths) else 1
        padded = torch.full((len(seqs), maxl), PAD_IDX, dtype=torch.long)
        for i, s in enumerate(seqs):
            padded[i, :len(s)] = s

        labels = torch.tensor(labels, dtype=torch.float32) # binary
        return padded.to(DEVICE), lengths.to(DEVICE), labels.to(DEVICE)

    train_ds = TextDataset(train_texts, y_train)
    val_ds = TextDataset(val_texts, y_val)
    test_ds = TextDataset(test_texts, y_test)

    train_loader = DataLoader(train_ds, batch_size=BATCH_SIZE, shuffle=True, ↴
        collate_fn=collate_batch)
    val_loader = DataLoader(val_ds, batch_size=BATCH_SIZE, shuffle=False, ↴
        collate_fn=collate_batch)
    test_loader = DataLoader(test_ds, batch_size=BATCH_SIZE, shuffle=False, ↴
        collate_fn=collate_batch)

```

Using device: cuda
Vocab size: 36,117 (PAD=0, UNK=1)

```
[4]: # Model
class LSTMClassifier(nn.Module):
    def __init__(self, vocab_size, emb_dim, hid_dim, num_layers, bidir, ↴
        dropout, pad_idx):
        super().__init__()
        self.embedding = nn.Embedding(vocab_size, emb_dim, padding_idx=pad_idx)
        self.lstm = nn.LSTM(
            input_size=emb_dim,
            hidden_size=hid_dim,
            num_layers=num_layers,
            batch_first=True,
            bidirectional=bidir,
            dropout=0.0 if num_layers == 1 else dropout
        )
        out_dim = hid_dim * (2 if bidir else 1)
```

```

        self.dropout = nn.Dropout(dropout)
        self.fc = nn.Linear(out_dim, 1)

    def forward(self, x, lengths):
        # x: [B, T]
        emb = self.dropout(self.embedding(x)) # [B, T, E]

        # pack sequences for speed
        lengths_cpu = lengths.to("cpu")
        packed = nn.utils.rnn.pack_padded_sequence(emb, lengths_cpu, □
        ↵batch_first=True, enforce_sorted=False)
        packed_out, (h, c) = self.lstm(packed)

        # h shape: [num_layers * num_directions, B, H]
        if self.lstm.bidirectional:
            # last layer forward + backward
            h_f = h[-2, :, :]
            h_b = h[-1, :, :]
            h_cat = torch.cat([h_f, h_b], dim=1) # [B, 2H]
        else:
            h_cat = h[-1, :, :] # [B, H]

        logits = self.fc(self.dropout(h_cat)).squeeze(1) # [B]
        return logits

model = LSTMClassifier(
    vocab_size=len(itos),
    emb_dim=EMB_DIM,
    hid_dim=HID_DIM,
    num_layers=NUM_LAYERS,
    bidir=BIDIR,
    dropout=DROPOUT,
    pad_idx=PAD_IDX
).to(DEVICE)

criterion = nn.BCEWithLogitsLoss()
optimizer = torch.optim.AdamW(model.parameters(), lr=LR)

```

[5]:

```

# -----
# Metrics helpers
# -----
@torch.no_grad()
def predict_loader(loader):
    model.eval()
    all_probs, all_y = [], []
    for x, lengths, y in loader:
        logits = model(x, lengths)

```

```

        probs = torch.sigmoid(logits).detach().cpu().numpy()
        all_probs.append(probs)
        all_y.append(y.detach().cpu().numpy())
    probs = np.concatenate(all_probs) if all_probs else np.array([])
    ytrue = np.concatenate(all_y).astype(int) if all_y else np.array([],  

dtype=int)
    return probs, ytrue

def eval_split(loader, threshold=0.5):
    probs, ytrue = predict_loader(loader)
    ypred = (probs >= threshold).astype(int)
    acc = accuracy_score(ytrue, ypred)
    p, r, f1, _ = precision_recall_fscore_support(ytrue, ypred,  

average="binary", zero_division=0)
    return acc, p, r, f1

```

```
[6]: # -----
# Train loop with timing + early stopping
# -----
best_val_f1 = -1.0
best_state = None
no_improve = 0
epoch_times = []

for epoch in range(1, EPOCHS + 1):
    train_loss_hist = []
    val_f1_hist = []
    val_acc_hist = []
    t0 = time.time()
    model.train()
    running_loss = 0.0
    n_batches = 0

    for x, lengths, y in train_loader:
        optimizer.zero_grad()
        logits = model(x, lengths)
        loss = criterion(logits, y)
        loss.backward()
        nn.utils.clip_grad_norm_(model.parameters(), CLIP)
        optimizer.step()

        running_loss += float(loss.item())
        n_batches += 1

    train_loss = running_loss / max(n_batches, 1)
    val_acc, val_p, val_r, val_f1 = eval_split(val_loader)
```

```

t1 = time.time()
epoch_sec = t1 - t0
epoch_times.append(epoch_sec)
train_loss_hist.append(train_loss)
val_f1_hist.append(val_f1)
val_acc_hist.append(val_acc)

# After epoch 1, print a rough projection based on observed time/epoch
if epoch == 1:
    projected = epoch_sec * EPOCHS
    print(f"\n[Timing] Epoch 1 took {epoch_sec:.1f}s. Rough projected total for {EPOCHS} epochs: ~{projected/60:.1f} min (before early stopping.)")

    print(f"Epoch {epoch}/{EPOCHS} | train_loss={train_loss:.4f} | "
          f"VAL acc={val_acc:.4f} p={val_p:.4f} r={val_r:.4f} f1={val_f1:.4f} | "
          f"time={epoch_sec:.1f}s")

# Early stopping on VAL F1
if val_f1 > best_val_f1 + 1e-4:
    best_val_f1 = val_f1
    best_state = {k: v.cpu().clone() for k, v in model.state_dict().items()}
    no_improve = 0
else:
    no_improve += 1
    if no_improve >= PATIENCE:
        print(f"Early stopping: no VAL F1 improvement for {PATIENCE} epochs.")
        break

# Restore best model
if best_state is not None:
    model.load_state_dict(best_state)

```

[Timing] Epoch 1 took 40.0s. Rough projected total for 8 epochs: ~5.3 min (before early stopping).

Epoch 1/8 | train_loss=0.1840 | VAL acc=0.9715 p=0.9675 r=0.9744 f1=0.9710 | time=40.0s

Epoch 2/8 | train_loss=0.0566 | VAL acc=0.9842 p=0.9931 r=0.9744 f1=0.9837 | time=37.8s

Epoch 3/8 | train_loss=0.0293 | VAL acc=0.9858 p=0.9956 r=0.9751 f1=0.9853 | time=37.6s

Epoch 4/8 | train_loss=0.0161 | VAL acc=0.9906 p=0.9922 r=0.9886 f1=0.9904 | time=37.5s

Epoch 5/8 | train_loss=0.0084 | VAL acc=0.9851 p=0.9985 r=0.9708 f1=0.9845 |

```

time=37.4s
Epoch 6/8 | train_loss=0.0092 | VAL acc=0.9913 p=0.9953 r=0.9868 f1=0.9911 |
time=37.4s
Epoch 7/8 | train_loss=0.0045 | VAL acc=0.9913 p=0.9936 r=0.9886 f1=0.9911 |
time=37.5s
Epoch 8/8 | train_loss=0.0039 | VAL acc=0.9918 p=0.9925 r=0.9908 f1=0.9916 |
time=37.6s

```

[7]:

```

# -----
# Final evaluation
# -----
val_acc, val_p, val_r, val_f1 = eval_split(val_loader)
test_acc, test_p, test_r, test_f1 = eval_split(test_loader)

print("\n===== FINAL (best checkpoint) =====")
print(f"VAL acc={val_acc:.4f} p={val_p:.4f} r={val_r:.4f} f1={val_f1:.4f}")
print(f"TEST acc={test_acc:.4f} p={test_p:.4f} r={test_r:.4f} f1={test_f1:.4f}")
print(f"Avg time/epoch: {np.mean(epoch_times):.1f}s over {len(epoch_times)}\u2192 epoch(s)")

```

```

===== FINAL (best checkpoint) =====
VAL acc=0.9918 p=0.9925 r=0.9908 f1=0.9916
TEST acc=0.7109 p=0.8016 r=0.8279 f1=0.8145
Avg time/epoch: 37.8s over 8 epoch(s)

```

[9]:

```

import numpy as np
import pandas as pd
import matplotlib.pyplot as plt

from sklearn.metrics import (
    confusion_matrix, ConfusionMatrixDisplay,
    precision_recall_fscore_support, accuracy_score,
    roc_auc_score, average_precision_score
)

# -----
# 1) Evaluate on a loader (return probs + labels)
# -----
@torch.no_grad()
def lstm_probs_on_loader(loader):
    model.eval()
    all_probs, all_y = [], []
    for x, lengths, y in loader:
        logits = model(x, lengths)
        probs = torch.sigmoid(logits).detach().cpu().numpy()
        all_probs.append(probs)

```

```

        all_y.append(y.detach().cpu().numpy())
probs = np.concatenate(all_probs) if all_probs else np.array([])
ytrue = np.concatenate(all_y).astype(int) if all_y else np.array([], dtype=int)
return probs, ytrue

def metrics_from_probs(ytrue, probs, thr=0.5):
    ypred = (probs >= thr).astype(int)
    acc = accuracy_score(ytrue, ypred)
    p, r, f1, _ = precision_recall_fscore_support(ytrue, ypred, average="binary", zero_division=0)

    out = {"acc": acc, "prec": p, "rec": r, "f1": f1}
    try: out["roc_auc"] = roc_auc_score(ytrue, probs)
    except Exception: out["roc_auc"] = np.nan
    try: out["pr_auc"] = average_precision_score(ytrue, probs)
    except Exception: out["pr_auc"] = np.nan
    return out, ypred

# -----
# 2) Plots
# -----
def plot_training_curves(train_loss_hist, val_f1_hist):
    epochs = np.arange(1, len(train_loss_hist) + 1)

    fig, ax = plt.subplots()
    ax.plot(epochs, train_loss_hist, marker="o")
    ax.set_title("LSTM Training Loss")
    ax.set_xlabel("Epoch")
    ax.set_ylabel("Loss")
    plt.show()

    fig, ax = plt.subplots()
    ax.plot(epochs, val_f1_hist, marker="o")
    ax.set_title("LSTM Validation F1")
    ax.set_xlabel("Epoch")
    ax.set_ylabel("F1")
    ax.set_ylim(0, 1)
    plt.show()

def plot_confidence_distributions(ytrue, probs, marker_prob=None):
    p_h = probs[ytrue == 0]
    p_a = probs[ytrue == 1]
    bins = np.linspace(0, 1, 51)

    fig, ax = plt.subplots()
    ax.hist(p_h, bins=bins, alpha=0.6, density=True, label="True Human (0)")

```

```

ax.hist(p_a, bins=bins, alpha=0.6, density=True, label="True AI (1)")
if marker_prob is not None:
    ax.axvline(marker_prob, linestyle="--", linewidth=2)
    ax.text(marker_prob, ax.get_ylim()[1]*0.95, "Custom input", ↴
rotation=90, va="top")
ax.set_title("LSTM Prob(AI) distributions (SCRAPED TEST)")
ax.set_xlabel("Prob(AI)")
ax.set_ylabel("Density")
ax.legend()
plt.show()

def plot_length_bin_f1(test_df, ytrue, probs, thr=0.5):
    lens = test_df["text"].astype(str).str.len()
    q1, q2 = lens.quantile([0.33, 0.66]).values
    bins = pd.cut(lens, [-np.inf, q1, q2, np.inf], ↴
labels=["short", "medium", "long"]).astype(str)

    ypred = (probs >= thr).astype(int)

    labels = ["short", "medium", "long"]
    f1s, counts = [], []
    for b in labels:
        m = (bins == b).values
        counts.append(int(m.sum()))
        if m.sum() == 0:
            f1s.append(np.nan)
        else:
            _, _, f1, _ = precision_recall_fscore_support(ytrue[m], ypred[m], ↴
average="binary", zero_division=0)
            f1s.append(float(f1))

    fig, ax = plt.subplots()
    ax.bar(labels, f1s)
    ax.set_xlim(0, 1)
    ax.set_xlabel("F1")
    ax.set_title("LSTM F1 by Length Bin (SCRAPED TEST)")
    for i, c in enumerate(counts):
        ax.text(i, 0.02, f'n={c}', ha="center", va="bottom")
    plt.show()

def per_source_table(test_df, ytrue, probs, thr=0.5, min_n=20):
    ypred = (probs >= thr).astype(int)
    rows = []
    for src, idx in test_df.groupby("source").indices.items():
        idx = np.array(list(idx))
        if len(idx) < min_n:
            continue

```

```

    m, _ = metrics_from_probs(ytrue[idx], probs[idx], thr=thr)
    rows.append({"source": src, "n": len(idx), **m})
out = pd.DataFrame(rows).sort_values("f1", ascending=False)
print(f"\n==== LSTM per-source performance (SCRAPED TEST, min_n={min_n})\n====")
display(out.round(4))
return out

def show_confident_errors(test_df, ytrue, probs, topk=10):
    df_err = test_df.copy().reset_index(drop=True)
    df_err["y_true"] = ytrue
    df_err["prob_ai"] = probs
    df_err["y_pred"] = (probs >= 0.5).astype(int)
    df_err["correct"] = (df_err["y_true"] == df_err["y_pred"])

    fp = df_err[(df_err["y_true"] == 0) & (df_err["y_pred"] == 1)].\
        sort_values("prob_ai", ascending=False).head(topk)
    fn = df_err[(df_err["y_true"] == 1) & (df_err["y_pred"] == 0)].\
        sort_values("prob_ai", ascending=True).head(topk)

    print(f"\n==== Top {topk} confident FALSE POSITIVES (Human predicted AI)\n====")
    display(fp[["source", "prob_ai", "text"]].assign(text=fp["text"].str.slice(0, 300) + "..."))

    print(f"\n==== Top {topk} confident FALSE NEGATIVES (AI predicted Human)\n====")
    display(fn[["source", "prob_ai", "text"]].assign(text=fn["text"].str.slice(0, 300) + "..."))


# -----
# 3) Run dashboard
# -----
# training curves (only if you logged these in Cell 1)
if "train_loss_hist" in globals() and "val_f1_hist" in globals():
    plot_training_curves(train_loss_hist, val_f1_hist)
else:
    print("No training curves found. Add train_loss_hist / val_f1_hist logging\nas suggested.")


# evaluate val + test
val_probs, val_ytrue = lstm_probs_on_loader(val_loader)
test_probs, test_ytrue = lstm_probs_on_loader(test_loader)

val_m, _ = metrics_from_probs(val_ytrue, val_probs, thr=0.5)
test_m, test_pred = metrics_from_probs(test_ytrue, test_probs, thr=0.5)

```

```
print("\n==== LSTM Compact Metrics (thr=0.5) ====")
display(pd.DataFrame([
    {"split":"VAL", **val_m},
    {"split":"TEST", **test_m},
]).round(4))
```






```
==== LSTM Compact Metrics (thr=0.5) ===
```


split	acc	prec	rec	f1	roc_auc	pr_auc
0 VAL	0.9918	0.9925	0.9908	0.9916	0.9994	0.9995
1 TEST	0.7109	0.8016	0.8279	0.8145	0.7460	0.9137


```
[10]: # confusion matrix on scraped test
cm = confusion_matrix(test_ytrue, test_pred, labels=[0,1])
disp = ConfusionMatrixDisplay(cm, display_labels=["Human(0)", "AI(1)"])
fig, ax = plt.subplots()
disp.plot(ax=ax, values_format="d")
ax.set_title("LSTM Confusion Matrix (SCRAPED TEST)")
plt.show()
```



```
[11]: # confidence distributions
plot_confidence_distributions(test_ytrue, test_probs)

# length-bin performance
plot_length_bin_f1(test_df.reset_index(drop=True), test_ytrue, test_probs, ↴
thr=0.5)
```





```
[18]: # per-source
per_source_table(test_df.reset_index(drop=True), test_ytrue, test_probs, thr=0.
                 ↪5, min_n=20)
```

```
==== LSTM per-source performance (SCRAPED TEST, min_n=20) ====
/usr/local/lib/python3.12/dist-packages/sklearn/metrics/_ranking.py:379:
UndefinedMetricWarning: Only one class is present in y_true. ROC AUC score is
not defined in that case.
    warnings.warn(
/usr/local/lib/python3.12/dist-packages/sklearn/metrics/_ranking.py:1033:
UserWarning: No positive class found in y_true, recall is set to one for all
thresholds.
    warnings.warn(
/usr/local/lib/python3.12/dist-packages/sklearn/metrics/_ranking.py:379:
UndefinedMetricWarning: Only one class is present in y_true. ROC AUC score is
not defined in that case.
    warnings.warn(
/usr/local/lib/python3.12/dist-packages/sklearn/metrics/_ranking.py:379:
UndefinedMetricWarning: Only one class is present in y_true. ROC AUC score is
not defined in that case.
```

```

warnings.warn(
/usr/local/lib/python3.12/dist-packages/sklearn/metrics/_ranking.py:1033:
UserWarning: No positive class found in y_true, recall is set to one for all
thresholds.

warnings.warn(
/usr/local/lib/python3.12/dist-packages/sklearn/metrics/_ranking.py:379:
UndefinedMetricWarning: Only one class is present in y_true. ROC AUC score is
not defined in that case.

warnings.warn(
/usr/local/lib/python3.12/dist-packages/sklearn/metrics/_ranking.py:1033:
UserWarning: No positive class found in y_true, recall is set to one for all
thresholds.

warnings.warn(

```

	source	n	acc	prec	rec	f1	roc_auc	pr_auc
1	chatgpt	1249	0.8279	1.0	0.8279	0.9058	NaN	1.0
0	arxiv	62	0.0806	0.0	0.0000	0.0000	NaN	0.0
2	news	97	0.4948	0.0	0.0000	0.0000	NaN	0.0
3	wikipedia	221	0.3213	0.0	0.0000	0.0000	NaN	0.0

[18]:

	source	n	acc	prec	rec	f1	roc_auc	pr_auc
1	chatgpt	1249	0.827862	1.0	0.827862	0.905826	NaN	1.0
0	arxiv	62	0.080645	0.0	0.000000	0.000000	NaN	0.0
2	news	97	0.494845	0.0	0.000000	0.000000	NaN	0.0
3	wikipedia	221	0.321267	0.0	0.000000	0.000000	NaN	0.0

[19]:

```

# confident errors
show_confident_errors(test_df.reset_index(drop=True), test_ytrue, test_probs, ↴
                      topk=10)

# expose test_probs for custom input overlay later
LSTM_TEST_PROBS = test_probs
LSTM_TEST_YTRUE = test_ytrue

```

==== Top 10 confident FALSE POSITIVES (Human predicted AI) ===

	source	prob_ai	text
115	arxiv	1.0	An Information Theoretic approach to Post Rand...
350	arxiv	1.0	Ray dynamics of whispering-gallery modes in cu...
831	arxiv	1.0	Unsupervised Learning of Solutions to Differen...
643	arxiv	1.0	Computing Valid p-values for Image Segmentatio...
692	wikipedia	1.0	late 20th and early 21st centuries has led to ...
776	wikipedia	1.0	four main ethical dimensions, defined as follo...
1075	arxiv	1.0	Conditional Independence Beyond Domain Separab...
935	arxiv	1.0	Estimation of seasonal long-memory parameters...
1309	arxiv	1.0	we provide long overlooked evidence that (3) e...
1378	wikipedia	1.0	constant rates without the fluctuations in une...

==== Top 10 confident FALSE NEGATIVES (AI predicted Human) ===

	source	prob_ai	text
398	chatgpt	2.296798e-08	Train min/max/avg: {tr_min} / {tr_max} / {tr_a...
213	chatgpt	2.696999e-08	filter={ "\$and": [{"payer_id": payer_id.lower...
668	chatgpt	3.472035e-08	using the same function as training feats = co...
1000	chatgpt	4.067079e-08	what the agent prioritizes/filters) agent_urls...
1221	chatgpt	5.867082e-08	scripts/<set>/__main__.py candidate = cur / "s...
453	chatgpt	6.216964e-08	== "0": prev_lab = "0"; continue pref, _, lab ...
1611	chatgpt	8.376617e-08	/ image_features.norm(dim=-1, keepdim=True) sim...
427	chatgpt	8.626050e-08	advances to Step 2) # =====
1144	chatgpt	9.531081e-08	r.rxn_code as oac_code, 'PH-' r.rxn_code...
1240	chatgpt	1.005773e-07	Step 3.") st.rerun() if back_to_1: st.session_...

[]:

[]:

[]:

[]:

[]:

[]:

[]:

[]:

[]:

[]:

[]:

[20]: # from https://gist.github.com/jonathanagustin/b67b97ef12c53a8dec27b343dca4abba
install can take a minute

```
import os
# @title Convert Notebook to PDF. Save Notebook to given directory
NOTEBOOKS_DIR = "/content/drive/MyDrive/" # @param {type:"string"}
NOTEBOOK_NAME = "04_lstm_sequence_model.ipynb" # @param {type:"string"}
#-----
from google.colab import drive
drive.mount("/content/drive/", force_remount=True)
NOTEBOOK_PATH = f"{NOTEBOOKS_DIR}/{NOTEBOOK_NAME}"
assert os.path.exists(NOTEBOOK_PATH), f"NOTEBOOK NOT FOUND: {NOTEBOOK_PATH}"
```

```

!apt install -y texlive-xetex texlive-fonts-recommended texlive-plain-generic > /dev/null 2>&1
!apt install pandoc > /dev/null 2>&1
!jupyter nbconvert "$NOTEBOOK_PATH" --to pdf > /dev/null 2>&1
NOTEBOOK_PDF = NOTEBOOK_PATH.rsplit('.', 1)[0] + '.pdf'
assert os.path.exists(NOTEBOOK_PDF), f"ERROR MAKING PDF: {NOTEBOOK_PDF}"
print(f"PDF CREATED: {NOTEBOOK_PDF}")

```

KeyboardInterrupt Traceback (most recent call last)

/tmp/ipython-input-4169629341.py in <cell line: 0>()

8 ↴

↳#-----#

9 from google.colab import drive

---> 10 drive.mount("/content/drive/", force_remount=True)

11 NOTEBOOK_PATH = f"{NOTEBOOKS_DIR}/{NOTEBOOK_NAME}"

12 assert os.path.exists(NOTEBOOK_PATH), f"NOTEBOOK NOT FOUND: ↴

↳{NOTEBOOK_PATH}"

/usr/local/lib/python3.12/dist-packages/google/colab/drive.py in ↴

↳mount(mountpoint, force_remount, timeout_ms, readonly)

95 def mount(mountpoint, force_remount=False, timeout_ms=120000, ↴

↳readonly=False):

96 """Mount your Google Drive at the specified mountpoint path."""

---> 97 return _mount(

98 mountpoint,

99 force_remount=force_remount,

/usr/local/lib/python3.12/dist-packages/google/colab/drive.py in ↴

↳_mount(mountpoint, force_remount, timeout_ms, ephemeral, readonly)

250

251 while True:

--> 252 case = d.expect([

253 success,

254 prompt,

/usr/local/lib/python3.12/dist-packages/pexpect/spawnbase.py in expect(self, ↴

↳pattern, timeout, searchwindowsize, async_, **kw)

352

353 compiled_pattern_list = self.compile_pattern_list(pattern)

--> 354 return self.expect_list(compiled_pattern_list,

355 timeout, searchwindowsize, async_)

356

/usr/local/lib/python3.12/dist-packages/pexpect/spawnbase.py in ↴

↳expect_list(self, pattern_list, timeout, searchwindowsize, async_, **kw)

```
381             return expect_async(exp, timeout)
382     else:
--> 383         return exp.expect_loop(timeout)
384
385     def expect_exact(self, pattern_list, timeout=-1, searchwindowsize=- ,

```

/usr/local/lib/python3.12/dist-packages/pexpect/expect.py in expect_loop(self, timeout)

```
    incoming = spawn.read_nonblocking(spawn.maxread, timeout)
    if self.spawn.delayafterread is not None:
--> 171        time.sleep(self.spawn.delayafterread)
    idx = self.new_data(incoming)
    # Keep reading until exception or return.
```

KeyboardInterrupt: